

```

"""Focusing on Anger & Surprise"""
from simpletransformers.classification import ClassificationModel,
ClassificationArgs
from sklearn.metrics import classification_report
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

```

## Models

```

model_args = ClassificationArgs( # args
    num_train_epochs=3, # cycle through data 6 times
    overwrite_output_dir=True, # o-write direct
    learning_rate=2e-5, # adjust rate
    train_batch_size=8, #training batches
    eval_batch_size=8,
    evaluate_during_training=True, # eval during train
)

#considered as the optim.
model_distilbert = ClassificationModel('distilbert', 'distilbert-base-
uncased', args=model_args , num_labels = 2)

```

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight', 'pre\_classifier.bias', 'pre\_classifier.weight']  
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```

model_roberta = ClassificationModel('roberta', 'roberta-base',
args=model_args, num_labels = 2)

```

Some weights of RobertaForSequenceClassification were not initialized from the model checkpoint at roberta-base and are newly initialized: ['classifier.dense.bias', 'classifier.dense.weight', 'classifier.out\_proj.bias', 'classifier.out\_proj.weight']  
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

## Dataset (1)

[https://colab.research.google.com/drive/19kBHJfud62yDHubqunrpbUxn\\_Qg8nJKZ?usp=sharing#scrollTo=ctfhqcQzJEXM](https://colab.research.google.com/drive/19kBHJfud62yDHubqunrpbUxn_Qg8nJKZ?usp=sharing#scrollTo=ctfhqcQzJEXM)

```

emo_train_df = pd.read_csv('C:\\Users\\NashW\\PycharmProjects\\
NeuralNetwork\\text_Mining_NLP_FGCU\\Datasets\\github-train.csv',
sep=',')
emo_test_df = pd.read_csv('C:\\Users\\NashW\\PycharmProjects\\
NeuralNetwork\\text_Mining_NLP_FGCU\\Datasets\\github-test.csv',
sep=',')

# preprocess (text and labels)
train_data = emo_train_df[['modified_comment', 'Anger', 'Surprise']]
test_data = emo_test_df[['modified_comment', 'Anger', 'Surprise']]

# setting our training splits
train_data, eval_data = (train_data.iloc[:int(0.8*len(train_data))],
                        train_data.iloc[int(0.8*len(train_data)):])

# Initial training of our first model
model_distilbert.train_model(train_data, eval_df=eval_data)

C:\\Users\\NashW\\PycharmProjects\\NeuralNetwork\\cuda\\lib\\site-packages\\
simpletransformers\\classification\\classification_model.py:610:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
  warnings.warn(

{"model_id": "763f98af167d4045a8373bb2117c5572", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "07e3794ab68d456aa1db0d7133bc4810", "version_major": 2, "vers
ion_minor": 0}

C:\\Users\\NashW\\PycharmProjects\\NeuralNetwork\\cuda\\lib\\site-packages\\
simpletransformers\\classification\\classification_model.py:882:
FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated.
Please use `torch.amp.GradScaler('cuda', args...)` instead.
  scaler = amp.GradScaler()

{"model_id": "4fb43a9b7403406b9227880a2a1b953e", "version_major": 2, "vers
ion_minor": 0}

C:\\Users\\NashW\\PycharmProjects\\NeuralNetwork\\cuda\\lib\\site-packages\\
simpletransformers\\classification\\classification_model.py:905:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
  with amp.autocast():
C:\\Users\\NashW\\PycharmProjects\\NeuralNetwork\\cuda\\lib\\site-packages\\
simpletransformers\\classification\\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
  warnings.warn(

```

```
{"model_id": "84f1397c067a42a48621cdd69841f46b", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1188: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
    warnings.warn(
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1033: UserWarning: No positive class found
in y_true, recall is set to one for all thresholds.
    warnings.warn(
```

```
{"model_id": "10d125ba53c94609a876b49038272517", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
    warnings.warn(
```

```
{"model_id": "3d5b483c99df4e63b57ea28d756cb4e2", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1188: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
    warnings.warn(
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1033: UserWarning: No positive class found
in y_true, recall is set to one for all thresholds.
    warnings.warn(
```

```
{"model_id": "12e5fbc035b84008afac4826f7234f33", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
    with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
```

```
    warnings.warn(
```

```
{"model_id": "df262719ba51465b969e6416f6674968", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
    with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1188: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
```

```
    warnings.warn(
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1033: UserWarning: No positive class found
in y_true, recall is set to one for all thresholds.
```

```
    warnings.warn(
```

```
(480,
 defaultdict(list,
              {'global_step': [160, 320, 480],
               'train_loss': [0.6623420715332031,
                              0.4143962860107422,
                              0.31427574157714844],
               'mcc': [0.0, 0.0, 0.0],
               'accuracy': [0.93125, 0.809375, 0.83125],
               'f1_score': [0.48220064724919093,
                             0.4473229706390328,
                             0.4539249146757679],
               'tp': [0, 0, 0],
               'tn': [298, 259, 266],
               'fp': [22, 61, 54],
               'fn': [0, 0, 0],
               'auROC': [nan, nan, nan],
               'auprc': [0.0, 0.0, 0.0],
               'eval_loss': [0.3437260627746582,
                              0.35400094985961916,
                              0.3017214298248291]}))
```

```

#evaluate the model based on the test data
result, model_output, incorrect_predictions =
model_distilbert.eval_model(test_data)

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
  warnings.warn(

{"model_id": "4f33755c51524451969ee0ed94adff46", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "5e5e3c67a5874cb3adb072d69cc2f220", "version_major": 2, "vers
ion_minor": 0}

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
  with amp.autocast():

# predict emotions
predictions =
model_distilbert.predict(test_data['modified_comment'].tolist())

{"model_id": "0e411073b29f4097ab15ee9440986fcf", "version_major": 2, "vers
ion_minor": 0}

{"model_id": "f39a39fe2d63415b8a43f506ecba96c9", "version_major": 2, "vers
ion_minor": 0}

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:2188:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
  with amp.autocast():

y_test_anger = test_data['Anger'].to_numpy()
y_test_surprise = test_data['Surprise'].to_numpy()

print(classification_report(y_test_anger, predictions[0]))

```

	precision	recall	f1-score	support
0	0.89	0.88	0.89	332
1	0.45	0.47	0.46	68
accuracy			0.81	400
macro avg	0.67	0.68	0.67	400
weighted avg	0.82	0.81	0.81	400

```
print(classification_report(y_test_surprise, predictions[0]))
```

	precision	recall	f1-score	support
0	0.85	0.83	0.84	335
1	0.21	0.23	0.22	65
accuracy			0.73	400
macro avg	0.53	0.53	0.53	400
weighted avg	0.74	0.73	0.74	400

```
# Roberta Model
```

```
model_roberta.train_model(train_data, eval_df=eval_data)
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:610:
```

```
UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "caeefefde4984fa082d6e263c6a98e0e", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "fa3cd005184f435cad69016e7562327b", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:882:
```

```
FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated. Please use `torch.amp.GradScaler('cuda', args...)` instead.
```

```
scaler = amp.GradScaler()
```

```
{"model_id": "903bd25cc837491c84908503298122e3", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:905:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1453:
```

```
UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "00d0163fea5a4a42a776b7dcfa5ec439", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1505:
```

```

FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_classification.py:407: UserWarning: A single label
was found in 'y_true' and 'y_pred'. For the confusion matrix to have
the correct shape, use the 'labels' parameter to pass all known
labels.
    warnings.warn(
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1188: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
    warnings.warn(
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1033: UserWarning: No positive class found
in y_true, recall is set to one for all thresholds.
    warnings.warn(

{"model_id": "525392035b6c4421ad5c554486fccfc5", "version_major": 2, "vers
ion_minor": 0}

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
    warnings.warn(

{"model_id": "667d6a89b5f341dd92b1e239dc099462", "version_major": 2, "vers
ion_minor": 0}

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1188: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
    warnings.warn(
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
sklearn\metrics\_ranking.py:1033: UserWarning: No positive class found
in y_true, recall is set to one for all thresholds.
    warnings.warn(

```



```
{"model_id": "36ccac2b5e1d4f158e6341006d0c1c01", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:905:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
```

```
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
    with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1453:
```

```
UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
```

```
    warnings.warn(
```

```
{"model_id": "a94e65a027804c1ebbbe2f3e89d2030d", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1505:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
```

```
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
    with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\sklearn\metrics\_ranking.py:1188: UndefinedMetricWarning: No positive samples in y_true, true positive value should be meaningless
```

```
    warnings.warn(
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\sklearn\metrics\_ranking.py:1033: UserWarning: No positive class found in y_true, recall is set to one for all thresholds.
```

```
    warnings.warn(
```

```
(480,
```

```
    defaultdict(list,
```

```
        {'global_step': [160, 320, 480],
```

```
         'train_loss': [0.38518333435058594,
```

```
                     0.26589345932006836,
```

```
                     0.09686851501464844],
```

```
         'mcc': [0.0, 0.0, 0.0],
```

```
         'accuracy': [1.0, 0.90625, 0.84375],
```

```
         'f1_score': [1.0, 0.47540983606557374,
```

```
0.4576271186440678],
```

```
         'tp': [0, 0, 0],
```

```
         'tn': [320, 290, 270],
```

```
         'fp': [0, 30, 50],
```

```
         'fn': [0, 0, 0],
```

```
         'auROC': [nan, nan, nan],
```

```
         'auprc': [0.0, 0.0, 0.0],
```

```
         'eval_loss': [0.08381292819976807,
```

```
                     0.21228610277175902,
```

```
                     0.2678683876991272]})
```

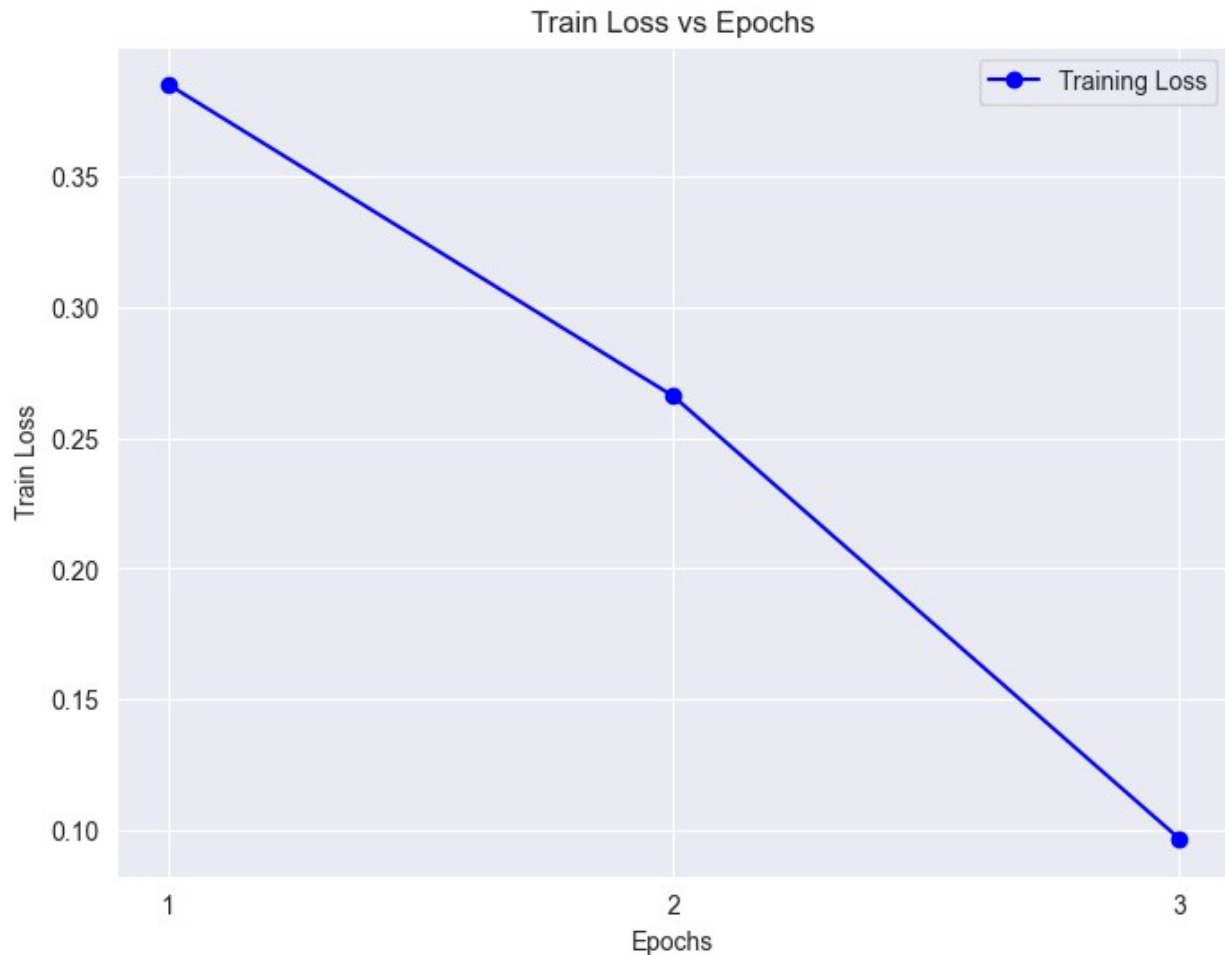


```
#train_loss': [0.43309783935546875, 0.8048801422119141,
0.042788028717041016]
# Define the epochs and corresponding eval_loss values
epochs = [1, 2, 3]
train_loss = [0.38518333435058594,
              0.26589345932006836,
              0.09686851501464844]

# Create the plot
plt.figure(figsize=(8, 6))
plt.plot(epochs, train_loss, marker='o', linestyle='-', color='b',
label='Training Loss')

# Add titles and labels
plt.title('Train Loss vs Epochs')
plt.xlabel('Epochs')
plt.ylabel('Train Loss')
plt.xticks(epochs) # Ensure the x-axis has ticks for each epoch

# Display the plot
plt.legend()
plt.grid(True)
plt.show()
```



*#evaluate the model based on the test data*

```
result, model_output, incorrect_predictions =  
model_roberta.eval_model(test_data)
```

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\  
simpletransformers\classification\classification\_model.py:1453:

UserWarning: Dataframe headers not specified. Falling back to using  
column 0 as text and column 1 as labels.

```
warnings.warn(
```

```
{"model_id": "b7e409ea936842c48e1bad4733fc59b8", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "e41632cbb9c3417082eeca17570f3ca6", "version_major": 2, "vers  
ion_minor": 0}
```

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\  
simpletransformers\classification\classification\_model.py:1505:

FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.

Please use `torch.amp.autocast('cuda', args...)` instead.

```
with amp.autocast():
```

```
# predict emotions
predictions =
model_roberta.predict(test_data['modified_comment'].tolist())

{"model_id": "139c2c8263b94721b756e04bb90814f3", "version_major": 2, "version_minor": 0}

{"model_id": "00a60857532a46f59435b1f6a70e8e59", "version_major": 2, "version_minor": 0}

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:2188:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
  with amp.autocast():

y_test_anger = test_data['Anger'].to_numpy()
y_test_surprise = test_data['Surprise'].to_numpy()

print(classification_report(y_test_anger, predictions[0]))
```

	precision	recall	f1-score	support
0	0.89	0.89	0.89	332
1	0.46	0.47	0.47	68
accuracy			0.82	400
macro avg	0.68	0.68	0.68	400
weighted avg	0.82	0.82	0.82	400

```
print(classification_report(y_test_surprise, predictions[0]))
```

	precision	recall	f1-score	support
0	0.83	0.82	0.83	335
1	0.14	0.15	0.15	65
accuracy			0.71	400
macro avg	0.49	0.49	0.49	400
weighted avg	0.72	0.71	0.72	400

## Dataset (2) Go Emotions

```
"""https://github.com/google-research/google-research/tree/master/
goemotions"""
```

```
go_emotions = pd.read_csv('https://raw.githubusercontent.com/google-
research/google-research/refs/heads/master/goemotions/data/train.tsv',
```

```

sep='\t')
go_emotions.columns = ["text", "targets", "id"]

#go_emotions["targets"] = go_emotions["targets"].to_numpy()
#go_emotions["text"] = go_emotions["text"].astype(str).tolist()

go_emotions["targets"].value_counts()
#Size of dataset = 43,410 rows

targets
27      12822
0       2710
4       1873
15      1857
1       1652
...
6,15,22      1
9,10,19      1
7,10,25      1
7,9,24,25    1
0,1,18       1
Name: count, Length: 711, dtype: int64

#Testing features
string_rows = go_emotions["text"].astype(str).tolist()
string_rows

#Testing labels
target_rows = go_emotions["targets"].astype(str).tolist()
target_rows

import re
t = []
for row in target_rows:
    t.extend(re.findall(r'\b2\b', row))

go_emotions["text"] = string_rows
go_emotions["text"]

0      Now if he does off himself, everyone will thin...
1      WHY THE FUCK IS BAYLESS ISOING
2      To make her feel threatened
3      Dirty Southern Wankers
4      OmG pEyToN iSn'T gOoD eNoUgH t0 hElP uS iN tHe...
...
43404  Added you mate well I've just got the bow and ...
43405  Always thought that was funny but is it a refe...
43406  What are you talking about? Anything bad that ...
43407      More like a baptism, with sexy results!
43408      Enjoy the ride!
Name: text, Length: 43409, dtype: object

```

```

#Library for cleaning our text

#Removing any user handles
import neattext.functions as nfx
go_emotions["Clean_txt"] =
go_emotions["text"].apply(nfx.remove_userhandles)

#Removing any stopwords
go_emotions["Clean_txt"] =
go_emotions["Clean_txt"].apply(nfx.remove_stopwords)

go_emotions["Clean_txt"]

0      himself, think hes laugh screwing people inste...
1                                     FUCK BAYLESS ISOING
2                                     feel threatened
3                                     Dirty Southern Wankers
4      OmG pEyToN g0oD hElP PLAyOfFs! Dumbass Broncos...
...
43404   Added mate I've got bow love hunting aspect ga...
43405                                     thought funny reference anything?
43406   talking about? bad happened [NAME] fault - goo...
43407                                     like baptism, sexy results!
43408                                     Enjoy ride!
Name: Clean_txt, Length: 43409, dtype: object

# 2, 14, 17, 18, 25, 26
# Those emotions covered in first dataset
# Anger(2), Surprise(26) == (Nash)
# Joy(17), Sadness(25) == (Isaiah)
# Love(18), Fear(14) == (Catalina)

#This is more of a reference as you go along
emotion_targets = {
    0 : "admiration",
    1 : "amusement",
    2 : "anger",
    3 : "annoyance",
    4 : "approval",
    5 : "caring",
    6 : "confusion",
    7 : "curiosity",
    8 : "desire",
    9 : "disappointment",
    10 : "disapproval",
    11 : "disgust",
    12 : "embarrassment",
    13 : "excitement",
    14 : "fear",
    15 : "gratitude",

```

```

16 : "grief",
17 : "joy",
18 : "love",
19 : "nervousness",
20 : "optimism",
21 : "pride",
22 : "realization",
23 : "relief",
24 : "remorse",
25 : "sadness",
26 : "surprise",
27 : "neutral"
}

# Specifically pulling anger and surprise, avoiding similar string-
like classes
import re

go_emotions["anger"] = [1 if re.search(r'\b2\b', str(row)) else 0 for
row in go_emotions["targets"]]

go_emotions["surprise"] = [1 if re.search(r'\b26\b', str(row)) else 0
for row in go_emotions["targets"]]

test = pd.concat([go_emotions["Clean_txt"], go_emotions["anger"],
go_emotions["surprise"]], axis=1)
go_emotions

```

	text	targets
id \		
0	Now if he does off himself, everyone will thin...	27
ed00q6i		
1	WHY THE FUCK IS BAYLESS ISOING	2
eezlygj		
2	To make her feel threatened	14
ed7ypvh		
3	Dirty Southern Wankers	3
ed0bdzj		
4	OmG pEyToN iSn'T gOoD eNoUgH tO hElP uS iN tHe...	26
edvnz26		
...	...	...
...		
43404	Added you mate well I've just got the bow and ...	18
edsb738		
43405	Always thought that was funny but is it a refe...	6
ee7fdou		
43406	What are you talking about? Anything bad that ...	3
efgbhks		
43407	More like a baptism, with sexy results!	13
edlnaf8		

```
43408          Enjoy the ride!      17
eecwmbq
```

```
          Clean_txt  anger
surprise
0    himself, think hes laugh screwing people inste...  0
0
1          FUCK BAYLESS ISOING      1
0
2          feel threatened          0
0
3    Dirty Southern Wankers        0
0
4    OmG pEyToN gOoD hElP PLAy0fFs! Dumbass Broncos...  0
1
...          ...      ...
...
43404  Added mate I've got bow love hunting aspect ga...  0
0
43405          thought funny reference anything?      0
0
43406  talking about? bad happened [NAME] fault - goo...  0
0
43407          like baptism, sexy results!            0
0
43408          Enjoy ride!              0
0
```

```
[43409 rows x 6 columns]
```

```
train_df_data = go_emotions[["Clean_txt", "anger", "surprise"]]
test_df_data = go_emotions[["Clean_txt", "anger", "surprise"]]
train_df_data[:10000]
```

```
          Clean_txt  anger
surprise
0    himself, think hes laugh screwing people inste...  0
0
1          FUCK BAYLESS ISOING      1
0
2          feel threatened          0
0
3    Dirty Southern Wankers        0
0
4    OmG pEyToN gOoD hElP PLAy0fFs! Dumbass Broncos...  0
1
...          ...      ...
..
9995  Meh, liga knew loved went away [NAME] fled Italy  0
0
```



```

9996 Meteor Man Edit: Hancock superhero movie terri...      0
0
9997 true.      0
0
9998 guy kept jumping bridge Eichenwald death over....      0
0
9999 Quick... there's boot licked today yet!      0
0

```

```
[10000 rows x 3 columns]
```

```
# use to eval during training
```

```
train_data, eval_data =
(train_df_data.iloc[ :int(0.8*len(train_df_data[0:10000]))],
```

```
train_df_data.iloc[int(0.8*len(train_df_data[0:10000])): ])
```

```
eval_data.shape
```

```
(35409, 3)
```

```
model_roberta.train_model(train_data, eval_df=eval_data)
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:610:
```

```
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "1404c625d18c46c1a127a3dd470fc968", "version_major": 2, "vers
ion_minor": 0}
```

```
{"model_id": "9570d2e67de6439991a68b8b98d9eeda", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:882:
```

```
FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated.
Please use `torch.amp.GradScaler('cuda', args...)` instead.
```

```
scaler = amp.GradScaler()
```

```
{"model_id": "366282548e0d4ff6a261eb4a218f4193", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
torch\optim\lr_scheduler.py:227: UserWarning: Detected call of
`lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and
```

later, you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step()`. Failure to do this will result in PyTorch skipping the first value of the learning rate schedule. See more details at <https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate>

```
warnings.warn(
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1453:
```

```
UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "63a0cd3165d04e98a9fa048a5cd73f09", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1505:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
{"model_id": "217252544139486e9d7054faac9e5f0f", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:905:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1453:
```

```
UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "7270196ce61a44a2b07e82f3f9b467cb", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1505:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:1453:
```

```
UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "61ce8f4d4d9c411892a914e9699a5d60", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
```

```
{"model_id": "4678607b32924feea063177c65f3597b", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
    warnings.warn(
```

```
{"model_id": "a48a59404d414d8994f710435bf90be1", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
```

```
(3000,
 defaultdict(list,
    {'global_step': [1000, 2000, 2000, 3000],
     'train_loss': [0.001577138900756836,
                    0.23310434818267822,
                    0.23310434818267822,
                    0.0013056397438049316],
     'mcc': [0.37432046086500753,
             0.389930802064812,
             0.389930802064812,
             0.387232420383696],
     'accuracy': [0.9634556186280324,
                  0.9628060662543421,
                  0.9628060662543421,
                  0.9592758903103731],
     'f1_score': [0.6813123819368836,
                  0.6917632375848333,
                  0.6917632375848333,
                  0.6932363843945294],
     'tp': [399, 444, 444, 496],
     'tn': [33716, 33648, 33648, 33471],
```

```
'fp': [406, 474, 474, 651],  
'fn': [888, 843, 843, 791],  
'auROC': [0.8916617105029273,  
0.886158444581163,  
0.886158444581163,  
0.8894676772731985],  
'auprc': [0.38169666380193523,  
0.3663623067441677,  
0.3663623067441677,  
0.3713972350016947],  
'eval_loss': [0.13536213350155996,  
0.17852038873461804,  
0.17852038873461804,  
0.19445478755274717]}}))
```

*# Define the epochs and corresponding train\_loss values from above*

```
epochs = [1, 2, 2, 3]  
train_loss = [0.001577138900756836,  
0.23310434818267822,  
0.23310434818267822,  
0.0013056397438049316]
```

*# Create the plot*


```
plt.figure(figsize=(8, 6))  
plt.plot(epochs, train_loss, marker='o', linestyle='-', color='b',  
label='Training Loss')
```

*# Add titles and labels*

```
plt.title('Train Loss vs Epochs')  
plt.xlabel('Epochs')  
plt.ylabel('Train Loss')  
plt.xticks(epochs) # Ensure the x-axis has ticks for each epoch
```

*# Display the plot*

```
plt.legend()  
plt.grid(True)  
plt.show()
```



```
#evaluate the model based on the test data
```

```
result, model_output, incorrect_predictions =  
model_roberta.eval_model(test_df_data)
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\  
simpletransformers\classification\classification_model.py:1453:
```

```
UserWarning: Dataframe headers not specified. Falling back to using  
column 0 as text and column 1 as labels.
```

```
warnings.warn(  

```

```
{"model_id": "1a273704bc054d4a99f2baaee5bb8e4e", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "d7819d71ca4649a0823744ae633ffc36", "version_major": 2, "vers  
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\  
simpletransformers\classification\classification_model.py:1505:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
```

```
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```

print(result)
{'mcc': 0.46346007824966506, 'accuracy': 0.9645465226105185,
 'f1_score': 0.7312954017398431, 'tp': 713, 'tn': 41157, 'fp': 685,
 'fn': 854, 'auroc': 0.9067187508531426, 'auprc': 0.4682701521976558,
 'eval_loss': 0.16857722570207737}

pred = model_roberta.predict(test_df_data["Clean_txt"].tolist())

{"model_id": "3ea3a17986784d44971786e6465e6867", "version_major": 2, "version_minor": 0}

{"model_id": "2588dc1911d34f96a5dd82e2cd5ebbbb6", "version_major": 2, "version_minor": 0}

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:2188:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
  with amp.autocast():

from sklearn.metrics import classification_report
y_test = test_df_data["anger"].to_numpy()
print(classification_report(y_test,
                           pred[0],
                           digits=4,
                           target_names=['Negative', 'Positive']))

```

	precision	recall	f1-score	support
Negative	0.9810	0.9840	0.9825	41842
Positive	0.5351	0.4914	0.5123	1567
accuracy			0.9662	43409
macro avg	0.7581	0.7377	0.7474	43409
weighted avg	0.9649	0.9662	0.9655	43409

```

from seaborn import heatmap
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.metrics import confusion_matrix

y_prediction_cm_1 = pred[0]
cm1 = confusion_matrix(y_test, y_prediction_cm_1)

```

```

## Accuracy Per Matrix is under each Graph
# Creating a heatmap visualization of the confusion matrix.
ax = sns.heatmap(cm1,
                  annot=True,
                  fmt='d',
                  cmap='crest',
                  linewidth=1,
                  xticklabels=['Negative [0]', 'Positive [1]'],
                  yticklabels=['Negative [0]', 'Positive [1]'])

## Configuring Orientation of labels
ax.xaxis.tick_top()
ax.set_yticklabels(ax.get_yticklabels(), rotation=0)

accuracy_cm1 = np.sum(np.diag(cm1)) / np.sum(cm1)

# Precision, Recall, and F1 Score Metrics
precision_cm1 = precision_score(y_test,
                                y_prediction_cm_1,
                                average='weighted',
                                zero_division=0)

recall_cm1 = recall_score(y_test,
                           y_prediction_cm_1,
                           average='weighted')

f1_cm1 = f1_score(y_test,
                  y_prediction_cm_1,
                  average='weighted')

plt.title(f"Confusion Matrix Heatmap\nRoberta Predicting Anger \
\n Accuracy-Score: {accuracy_cm1*100:.2f}%\nPrecision: \
{precision_cm1*100:.2f}%\n\
Recall: {recall_cm1*100:.2f}%\nF1: {f1_cm1*100:.2f}%\n")
plt.ylabel('<--True Label-->\n')
plt.xlabel('\n<--Predicted Label-->')


plt.show()

## The Heatmap is flipped, but you are absolutely allowed to do so
since our
## diagonal is correct with TP and TN

# 00 --> TN
# 01 --> FP
# 11 --> TP

```





```
y_test2 = test_df_data["surprise"].to_numpy()
print(classification_report(y_test2,
                             pred[0],
                             digits=4,
                             target_names=['Negative', 'Positive']))
```

	precision	recall	f1-score	support
Negative	0.9751	0.9664	0.9707	42349
Positive	0.0097	0.0132	0.0112	1060
accuracy			0.9431	43409
macro avg	0.4924	0.4898	0.4909	43409
weighted avg	0.9515	0.9431	0.9473	43409

```

y_prediction_cm_1 = pred[0]
cm1 = confusion_matrix(y_test2, y_prediction_cm_1)

## Accuracy Per Matrix is under each Graph
# Creating a heatmap visualization of the confusion matrix.
ax = sns.heatmap(cm1,
                  annot=True,
                  fmt='d',
                  cmap='crest',
                  linewidth=1,
                  xticklabels=['Negative [0]', 'Positive [1]'],
                  yticklabels=['Negative [0]', 'Positive [1]'])

## Configuring Orientation of labels
ax.xaxis.tick_top()
ax.set_yticklabels(ax.get_yticklabels(), rotation=0)

accuracy_cm1 = np.sum(np.diag(cm1)) / np.sum(cm1)

# Precision, Recall, and F1 Score Metrics
precision_cm1 = precision_score(y_test2,
                                y_prediction_cm_1,
                                average='weighted',
                                zero_division=0)

recall_cm1 = recall_score(y_test2,
                           y_prediction_cm_1,
                           average='weighted')

f1_cm1 = f1_score(y_test2,
                  y_prediction_cm_1,
                  average='weighted')

plt.title(f"Confusion Matrix Heatmap\nRoberta Predicting Surprised \
\n Accuracy-Score: {accuracy_cm1*100:.2f}%\nPrecision: \
{precision_cm1*100:.2f}%\n\
Recall: {recall_cm1*100:.2f}%\nF1: {f1_cm1*100:.2f}%\n")
plt.ylabel('<-- True Label-->\n')
plt.xlabel('\n<-- Predicted Label-->')


plt.show()

## The Heatmap is flipped, but you are absolutely allowed to do so
since our
## diagonal is correct with TP and TN

# 00 --> TN
# 01 --> FP

```

```
# 11 --> TP  
# 10 --> FN
```



```
model_distilbert.train_model(train_data, eval_df=eval_data)
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\simpletransformers\classification\classification_model.py:610:
```

```
UserWarning: DataFrame headers not specified. Falling back to using column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "04355ea5a07e478a9278c63d2fffa6f1", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0818538bafa348f6abde671267b40a3e", "version_major": 2, "version_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:882:
FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated.
Please use `torch.amp.GradScaler('cuda', args...)` instead.
    scaler = amp.GradScaler()
```

```
{"model_id": "e2abc7ea9b8741289ae194785df7bfc9", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
    warnings.warn(
```

```
{"model_id": "fb0dda55ccc947819902932926b0483a", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
```

```
{"model_id": "5f901e2a2d18475db3e6d4320a34047b", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
    warnings.warn(
```

```
{"model_id": "df0f1760681e43a1b7cf3c484fa4389e", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
    with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "341a1372c0564f92898a69baac36ac6e", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
```

```
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
{"model_id": "647527858597400d8e482efc8df5d8f0", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:905:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
```

```
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1453:
```

```
UserWarning: Dataframe headers not specified. Falling back to using
column 0 as text and column 1 as labels.
```

```
warnings.warn(
```

```
{"model_id": "6fbc84ac17e345fb9a7ea44f4664b9ce", "version_major": 2, "vers
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:1505:
```

```
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
```

```
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```

```
(3000,
 defaultdict(list,
   {'global_step': [1000, 2000, 2000, 3000],
    'train_loss': [0.001717209815979004,
                   0.000852733850479126,
                   0.000852733850479126,
                   0.0006740391254425049],
    'mcc': [0.32649626609026283,
            0.37144901727926743,
            0.37144901727926743,
            0.4020971964876041],
    'accuracy': [0.9659691039001383,
                 0.9667598633115875,
```

```

0.9667598633115875,
0.9596430286085459],
'f1_score': [0.6366627844305777,
0.6657944547662682,
0.6657944547662682,
0.7008585126982037],
'tp': [247, 315, 315, 523],
'tn': [33957, 33917, 33917, 33457],
'fp': [165, 205, 205, 665],
'fn': [1040, 972, 972, 764],
'auroc': [0.8975312179110315,
0.8763775641742937,
0.8763775641742937,
0.8897933973105416],
'auprc': [0.38366488681953465,
0.38598979789042,
0.38598979789042,
0.38628604535877253],
'eval_loss': [0.13695164171922872,
0.17672856645845844,
0.17672856645845844,
0.19368028894541411]})

```

*# Define the epochs and corresponding train\_loss values from above*

```

epochs = [1, 2, 2, 3]
train_loss = [0.001717209815979004,
0.000852733850479126,
0.000852733850479126,
0.0006740391254425049]

```

*# Create the plot*

```

plt.figure(figsize=(8, 6))
plt.plot(epochs, train_loss, marker='o', linestyle='-', color='b',
label='Training Loss')

```

*# Add titles and labels*

```

plt.title('Train Loss vs Epochs')
plt.xlabel('Epochs')
plt.ylabel('Train Loss')
plt.xticks(epochs) # Ensure the x-axis has ticks for each epoch

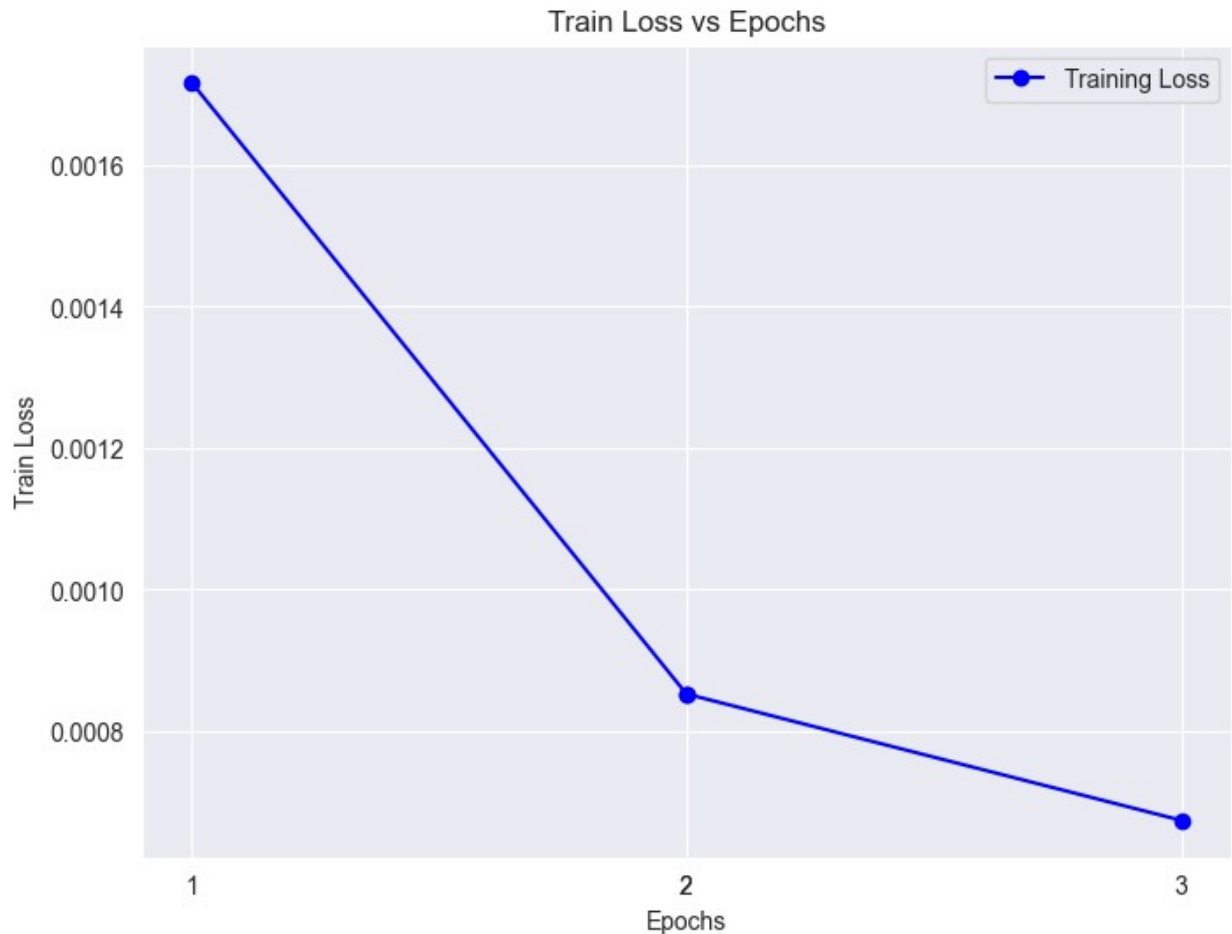
```

*# Display the plot*

```

plt.legend()
plt.grid(True)
plt.show()

```



```
#evaluate the model based on the test data
```

```
result, model_output, incorrect_predictions =  
model_distilbert.eval_model(test_df_data)
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\  
simpletransformers\classification\classification_model.py:1453:  
UserWarning: DataFrame headers not specified. Falling back to using  
column 0 as text and column 1 as labels.
```

```
warnings.warn(  

```

```
{"model_id": "e02dfefe2a75427cae2caf27cdaace2d", "version_major": 2, "vers  
ion_minor": 0}
```

```
{"model_id": "b4e6a7062a28426ea85068e3960bfd15", "version_major": 2, "vers  
ion_minor": 0}
```

```
C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\  
simpletransformers\classification\classification_model.py:1505:  
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.  
Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
with amp.autocast():
```



```

print(result)
{'mcc': 0.4953320639713798, 'accuracy': 0.9662282015250294,
 'f1_score': 0.7474085936210748, 'tp': 770, 'tn': 41173, 'fp': 669,
 'fn': 797, 'auroc': 0.9102141303625358, 'auprc': 0.5025116383047372,
 'eval_loss': 0.16170097528647076}

pred = model_distilbert.predict(test_df_data["Clean_txt"].tolist())

{"model_id": "24b1bb99889a4c988d0f9402528b0d3f", "version_major": 2, "version_minor": 0}

{"model_id": "425516e31c514dce92e5283fa236fea6", "version_major": 2, "version_minor": 0}

C:\Users\NashW\PycharmProjects\NeuralNetwork\cuda\lib\site-packages\
simpletransformers\classification\classification_model.py:2188:
FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated.
Please use `torch.amp.autocast('cuda', args...)` instead.
  with amp.autocast():

y_test3 = test_df_data["anger"].to_numpy()

print(classification_report(y_test3,
                           pred[0],
                           digits=4,
                           target_names=['Negative', 'Positive']))

```

	precision	recall	f1-score	support
Negative	0.9810	0.9840	0.9825	41842
Positive	0.5351	0.4914	0.5123	1567
accuracy			0.9662	43409
macro avg	0.7581	0.7377	0.7474	43409
weighted avg	0.9649	0.9662	0.9655	43409

```

y_prediction_cm_1 = pred[0]
cm1 = confusion_matrix(y_test3, y_prediction_cm_1)

## Accuracy Per Matrix is under each Graph
# Creating a heatmap visualization of the confusion matrix.
ax = sns.heatmap(cm1,
                 annot=True,
                 fmt='d',
                 cmap='crest',
                 linewidth=1,
                 xticklabels=['Negative [0]', 'Positive [1]'],
                 yticklabels=['Negative [0]', 'Positive [1]'])

## Configuring Orientation of labels

```

```

ax.xaxis.tick_top()
ax.set_yticklabels(ax.get_yticklabels(), rotation=0)

accuracy_cm1 = np.sum(np.diag(cm1)) / np.sum(cm1)

# Precision, Recall, and F1 Score Metrics
precision_cm1 = precision_score(y_test3,
                                y_prediction_cm_1,
                                average='weighted',
                                zero_division=0)

recall_cm1 = recall_score(y_test3,
                           y_prediction_cm_1,
                           average='weighted')

f1_cm1 = f1_score(y_test3,
                  y_prediction_cm_1,
                  average='weighted')

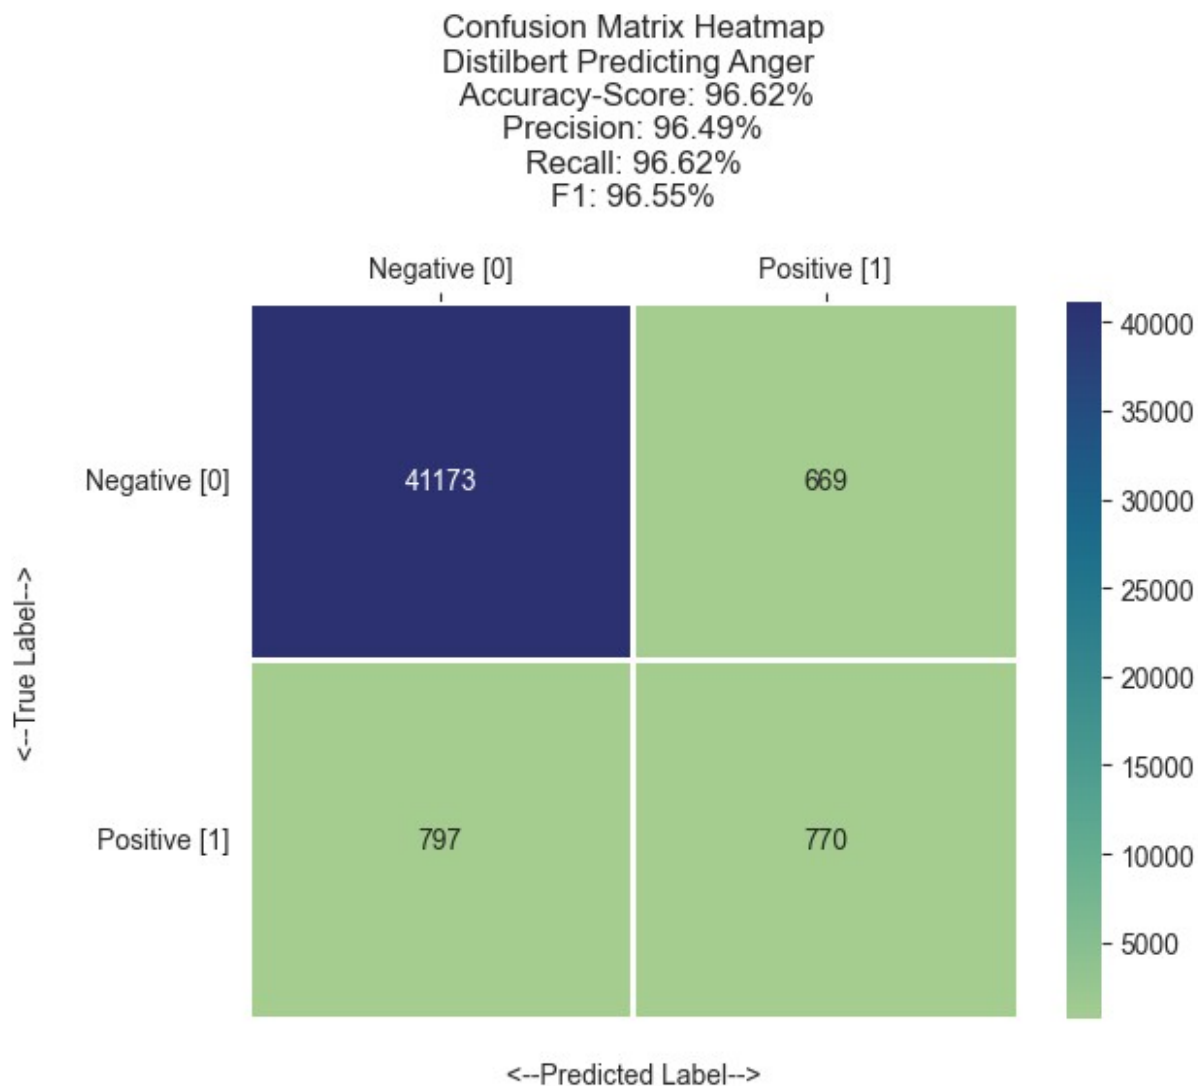
plt.title(f"Confusion Matrix Heatmap\nDistilbert Predicting Anger \
\n Accuracy-Score: {accuracy_cm1*100:.2f}%\nPrecision: \
{precision_cm1*100:.2f}%\n\
Recall: {recall_cm1*100:.2f}%\nF1: {f1_cm1*100:.2f}%\n")
plt.ylabel('<--True Label-->\n')
plt.xlabel('\n<--Predicted Label-->')

plt.show()

## The Heatmap is flipped, but you are absolutely allowed to do so
since our
## diagonal is correct with TP and TN

# 00 --> TN
# 01 --> FP
# 11 --> TP
# 10 --> FN

```



```
y_test4 = test_df_data["surprise"].to_numpy()
print(classification_report(y_test4,
                             pred[0],
                             digits=4,
                             target_names=[ 'Negative', 'Positive'])))
```

	precision	recall	f1-score	support
Negative	0.9751	0.9664	0.9707	42349
Positive	0.0097	0.0132	0.0112	1060
accuracy			0.9431	43409
macro avg	0.4924	0.4898	0.4909	43409
weighted avg	0.9515	0.9431	0.9473	43409

```

y_prediction_cm_1 = pred[0]
cm1 = confusion_matrix(y_test4, y_prediction_cm_1)

## Accuracy Per Matrix is under each Graph
# Creating a heatmap visualization of the confusion matrix.
ax = sns.heatmap(cm1,
                  annot=True,
                  fmt='d',
                  cmap='crest',
                  linewidth=1,
                  xticklabels=['Negative [0]', 'Positive [1]'],
                  yticklabels=['Negative [0]', 'Positive [1]'])

## Configuring Orientation of labels
ax.xaxis.tick_top()
ax.set_yticklabels(ax.get_yticklabels(), rotation=0)

accuracy_cm1 = np.sum(np.diag(cm1)) / np.sum(cm1)

# Precision, Recall, and F1 Score Metrics
precision_cm1 = precision_score(y_test4,
                                y_prediction_cm_1,
                                average='weighted',
                                zero_division=0)

recall_cm1 = recall_score(y_test4,
                           y_prediction_cm_1,
                           average='weighted')

f1_cm1 = f1_score(y_test4,
                  y_prediction_cm_1,
                  average='weighted')

plt.title(f"Confusion Matrix Heatmap\nDistilbert Predicting Surprised
\n
\n Accuracy-Score: {accuracy_cm1*100:.2f}%\nPrecision:
{precision_cm1*100:.2f}%\n\n
Recall: {recall_cm1*100:.2f}%\nF1: {f1_cm1*100:.2f}%\n")
plt.ylabel('<--True Label-->\n')
plt.xlabel('\n<--Predicted Label-->')

plt.show()

## The Heatmap is flipped, but you are absolutely allowed to do so
since our
## diagonal is correct with TP and TN

# 00 --> TN

```

```
# 01 --> FP
# 11 --> TP
# 10 --> FN
```

Confusion Matrix Heatmap  
Distilbert Predicting Surprised  
Accuracy-Score: 94.31%  
Precision: 95.15%  
Recall: 94.31%  
F1: 94.73%

